

## EC2 Dashboard

Events

Tags

Limits

### ▼ Instances

Instances **New**

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved

Instances **New**

Dedicated Hosts

Capacity Reservations

### ▼ Images

AMIs

### ▼ Elastic Block Store

Volumes

Snapshots

Lifecycle

Manager **New**

### ▼ Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

### ▼ Load Balancing

Load Balancers

Target Groups **New**

### ▼ Auto Scaling

Launch

Configurations

Auto Scaling Groups

## Events

Completed maintenance events are displayed on the Amazon EC2 console dashboard for up to a week. Actions for Instances Backed by Amazon EBS. You can wait for the maintenance to occur as scheduled. Alternatively, you can stop and start the instance, which migrates it to a new host.

## Tag

Each tag is a simple label consisting of a customer-defined key and an optional value that can make it easier to manage, search for, and filter resources by purpose, owner, environment, or other criteria. AWS tags can be used for many purposes. Tags for AWS Console Organization and Resource Groups.

## Limits

By default, AWS has a limit of 20 instances per region. This includes all instances set up on your AWS account. To increase EC2 limits, request a higher limit by providing information about the new limit and regions where it should be applied

## instances

EC2 virtual servers, also known as instances, are the building blocks for supercomputing on AWS. The EC2 dashboard contains resources and the ability to launch an instance. On the left-hand side of the dashboard there are many links, including to EC2 limits, instances, AMIs, Security Groups and SSH keys.

## Instance type

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instance types comprise varying combinations of CPU, memory, storage, and networking capacity and give you the flexibility to choose the appropriate mix of resources for your applications.

## launch template

A launch template is similar to a launch configuration, in that it specifies instance configuration information. Included are the ID of the Amazon Machine Image (AMI), the instance type, a key pair, security groups, and the other parameters that you use to launch EC2 instances.

## Spot Instance

A Spot Instance is an instance that uses spare EC2 capacity that is available for less than the On-Demand price. Because Spot Instances enable you to request unused EC2 instances at steep discounts, you can lower your Amazon EC2 costs significantly.

## Savings Plans

Savings Plans is a flexible pricing model offering lower prices compared to On-Demand pricing, in exchange for a specific usage commitment (measured in \$/hour) for a one or three-year period. ... The EC2 Instance Savings Plans apply to EC2 usage, and Amazon SageMaker Savings Plans apply to Amazon SageMaker usage.

## Reserved Instance (RI)

An Amazon Reserved Instance (RI) is a billing discount that allows you to save on your Amazon EC2 usage costs. When you purchase a Reserved Instance, you can set attributes such as instance type, platform, tenancy, Region, or Availability Zone (optional).

## Dedicated Host

An Amazon EC2 Dedicated Host is a physical server with EC2 instance capacity fully dedicated to your use. Dedicated Hosts can help you address compliance requirements and reduce costs by allowing you to use your existing server-bound software licenses.

## capacity reservation

On-Demand Capacity Reservations enable you to reserve compute capacity **for your Amazon EC2 instances in a specific Availability Zone for any duration**. This gives you the ability to create and manage Capacity Reservations independently from the billing discounts offered by Savings Plans or regional Reserved Instances.

## Amazon Machine Image (AMI)

An Amazon Machine Image (AMI) is a master image for the creation of virtual servers -- known as EC2 instances -- in the Amazon Web Services (AWS) environment. ... Block device mapping ensures that the correct volumes are attached to the launched instance.

## EBS volume

An Amazon EBS volume is a durable, block-level storage device that you can attach to your instances. ... EBS volumes persist independently from the running life of an EC2 instance. You can attach multiple EBS volumes to a single instance. The volume and instance must be in the same Availability Zone.

## Snapshot

A snapshot takes a copy of the EBS volume and places it in Amazon S3, where it is stored redundantly in multiple Availability Zones. The initial snapshot is a full copy of the volume; ongoing snapshots store incremental block-level changes only. This is a fast and reliable way to restore full volume data.

## Lifecycle Manager

With Amazon Data Lifecycle Manager, you can manage the lifecycle of your AWS resources. You create lifecycle policies, which are used to automate operations on the specified resources. Amazon DLM supports Amazon EBS volumes and snapshots.

## security group

A security group acts as a virtual firewall for your EC2 instances to control incoming and outgoing traffic. ... You can add rules to each security group that allow traffic to or from its associated instances. You can modify the rules for a security group at any time.

## Elastic IP

An Elastic IP provides you a single IP address that you can associate with different EC2 instances over time. If your EC2 instance has an Elastic IP and that instance is ever stopped or terminated, you can immediately associate a new EC2 instance with the Elastic IP.

## placement group

A cluster placement group is a logical grouping of instances within a single Availability Zone that benefit from low network latency, high network throughput. A spread placement group places instances on distinct hardware.

## key pair

A key pair, consisting of a public key and a private key, is a set of security credentials that you use to prove your identity when connecting to an EC2 instance. ... Anyone who possesses your private key can connect to your instances, so it's important that you store your private key in a secure place.

## network interface

An elastic network interface is a logical networking component in a VPC that represents a virtual network card. It can include the following attributes: A primary private IPv4 address from the IPv4 address range of your VPC. ... One or more IPv6 addresses. One or more security groups.

## Elastic Load Balancing

Elastic Load Balancing automatically distributes incoming application traffic across multiple targets, such as Amazon EC2 instances, containers, IP addresses, Lambda functions, and virtual appliances.

## target group

Each target group is used to route requests to one or more registered targets. When you create each listener rule, you specify a target group and conditions. When a rule condition is met, traffic is forwarded to the corresponding target group.

## Auto Scaling groups.

A launch configuration is an instance configuration template that an Auto Scaling group uses to launch EC2 instances. When you create a launch configuration, you specify information for the instances. ... You can specify your launch configuration with multiple Auto Scaling groups.

## Auto Scaling

When you enable Auto Scaling group metrics, your Auto Scaling group sends sampled data to CloudWatch every minute. There is no charge for enabling these metrics. To enable group metrics.